VIRTUE FOUNDATION INTELLIGENCE PLATFORM

Demo Recording Guide

Optimized for Twitter / LinkedIn | Target: 5 - 7 minutes

Databricks x Hack-Nation Hackathon | 2nd Place Winner

This guide provides the complete scene-by-scene recording flow for a tightened demo of the Virtue Foundation Intelligence Platform. All seven scenes from the original script are preserved but streamlined for pacing. Each scene includes voiceover script, on-screen actions, and timing targets. Record screen first, then layer voiceover in post-production.

Scene	Section	Duration
1	Hook - The Problem	0:00 - 0:35
2	IDP Agent (Data Pipeline)	0:35 - 1:40
3	Command Center (AI + Voice)	1:40 - 3:10
4	Strategic Planner	3:10 - 4:00
5	Facility Explorer	4:00 - 4:30
6	Data Integrity	4:30 - 4:50
7	Closing	4:50 - 5:10
	TOTAL TARGET	**5:10 - 6:00**

SCENE 1: HOOK - THE PROBLEM 0:00 - 0:35

VOICEOVER

"3.6 billion people lack access to essential healthcare. Not because resources don't exist - because the data is broken. Organizations like the Virtue Foundation deploy doctors and mobile clinics to underserved communities, but their facility data is scattered across messy spreadsheets and PDFs. A question like 'Which region has zero emergency care?' takes analysts weeks to answer. We built a platform that does it in seconds. This is the Virtue Foundation Intelligence Platform - second place at Databricks x Hack-Nation. Let me show you how it works."

ON-SCREEN ACTIONS

- Start on **Command Center** with Ghana map visible
- Slow pan across the map showing 797 facility markers
- Briefly hover over the **Regional Overview** panel and **ONLINE** status badge

TIP

Keep the mouse moving slowly. Let the visual speak. The map with hundreds of markers is the most shareable thumbnail frame for social media.

SCENE 2: IDP AGENT (DATA PIPELINE) 0:35 - 1:40

VOICEOVER

"Everything starts with the data. We took over 1,000 raw healthcare facility records from Ghana - unstructured, inconsistent, full of duplicates. Our Intelligent Document Processing agent uses GPT-4o-mini to read raw records and extract structured, usable data."

ON-SCREEN ACTIONS

- Click **IDP Agent** in the sidebar
- Click **View Schema** - show briefly, then close
- Select a facility from the dropdown
- Point out the **left panel** (raw CSV data) - highlight messy fields
- Point out the **right panel** - extracted capabilities with confidence scores

VOICEOVER

"Every extraction gets a confidence score. Surgery at 93% means strong evidence. Radiology at 58% means it needs human verification. We're not blindly trusting AI - we're giving decision-makers the information to know what to trust."

ON-SCREEN ACTIONS

- Click the **Agent Trace** tab - show step-by-step reasoning
- Click **Export JSON** - show clean structured output

VOICEOVER

"We processed 1,003 raw records into 797 clean, deduplicated, capability-scored profiles. And this pipeline works with any country's data - not just Ghana."

SCENE 3: COMMAND CENTER (AI + VOICE) 1:40 - 3:10

VOICEOVER

"Now that we have clean data - how do we make it useful? This is the Command Center."

ON-SCREEN ACTIONS

- Click **Command Center** in the sidebar
- Point out the map with facility markers and **medical desert zones** (red overlays)
- Point out the **Regional Health Overview** panel - stats, coverage %, key gaps

TEXT QUERY DEMO

ON-SCREEN ACTIONS

- Type: **"Which regions have no surgical capabilities?"**
- Wait for the response to generate

VOICEOVER

"I ask a plain English question that would normally take days of cross-referencing spreadsheets. Behind the scenes, a LangGraph multi-agent pipeline classifies the query, routes it to the right specialist agent, searches our FAISS vector store, and generates a cited answer."

ON-SCREEN ACTIONS

- Expand the **Agent Trace** - show step-by-step reasoning

- Point out **citations** linking back to facility records

VOICEOVER

"Every claim is cited. Full transparency - no black box."

VOICE QUERY DEMO

ON-SCREEN ACTIONS

- Click the **microphone button**
- Speak: **"What are the top medical deserts in Ghana?"**
- Wait for transcription + AI response
- Click **"Listen"** on the response - let it play 3-4 seconds
- Click **"Stop Audio"**

VOICEOVER

"Now imagine a field worker in a rural area - no keyboard, limited connectivity. They can speak directly. Speech-to-text with ElevenLabs Scribe, text-to-speech with ElevenLabs Multilingual v2. The platform is fully accessible without ever touching a keyboard."

TIP

This is the centerpiece demo. Rehearse the voice query 2-3 times before recording. Speak clearly at a normal pace. Test mic + TTS output volume beforehand.

SCENE 4: STRATEGIC PLANNER 3:10 - 4:00

VOICEOVER

"Finding gaps is one thing. Deciding what to do about them - that's the Strategic Planner."

ON-SCREEN ACTIONS

- Click **Strategic Planner** in the sidebar
- Point out the **stats cards** at the top - surgical deficit, capability gaps, maternal health gaps
- Show the **Medical Desert Matrix** - regions vs. capabilities (green = covered, red = desert)
- Point to specific red cells - **Upper West surgery**, **Savannah emergency care**

VOICEOVER

"10 healthcare capabilities across all 16 regions. Green means adequate coverage. Red means a medical desert - real communities with zero access."

ON-SCREEN ACTIONS

- Click/show **AI Recommendations**
- Briefly show **Deployment Plan Builder** - step through Scope and Assets

VOICEOVER

"The AI generates prioritized deployment recommendations - which region gets a surgical unit first and why. Population density, existing coverage, gap severity - all factored in. Decision-makers can build deployment plans directly from these recommendations. From data to action - in one platform."

SCENE 5: FACILITY EXPLORER 4:00 - 4:30

VOICEOVER

"When you need to go deep on a specific facility - search, filter, inspect."

ON-SCREEN ACTIONS

- Click **Facility Explorer** in the sidebar
- Type a facility name in search
- Apply a **region filter**

- Click **View Details** on a facility - show capabilities, confidence scores, raw source text
- Point out **anomaly detection** flag

VOICEOVER

"We flagged 43 anomalies across 39 facilities - like a small clinic claiming advanced surgery. You can't deploy resources based on data you don't trust. This feature builds that trust."

SCENE 6: DATA INTEGRITY 4:30 - 4:50

VOICEOVER

"The Data Integrity dashboard shows the full picture - field completeness rates, region-by-region data quality, and how our agents normalized messy region names into Ghana's 16 official regions."

ON-SCREEN ACTIONS

- Click **Data Integrity** in the sidebar
- Show stats cards - 797 unique facilities, completeness %, anomalies count
- Scroll to the **Region Name Normalization** table - show before/after corrections

TIP

Keep this scene brisk - 20 seconds max. The visuals carry the message. Don't linger on individual data points.

SCENE 7: CLOSING 4:50 - 5:10

ON-SCREEN ACTIONS

- Navigate back to **Command Center** - show the full Ghana map

VOICEOVER

"797 facilities. 10 medical deserts identified. 67 critical gaps. 43 data anomalies flagged. All in seconds - not weeks."

VOICEOVER

"Ghana is the proof of concept. The architecture is country-agnostic - upload facility data from any country and the entire intelligence layer rebuilds automatically."

VOICEOVER

"Built with React, FastAPI, LangGraph, GPT-4o-mini, FAISS, and ElevenLabs. Second place at Databricks x Hack-Nation. Every data point represents a patient who could receive care sooner."

PRE-RECORDING CHECKLIST

	Item	Details
	Backend	Running: `venv/bin/python -m uvicorn main:app --host 0.0.0.0 --port 8000`
	Frontend	Running: `npm run dev` on port 5173
	Chat	Clear all previous Command Center chat messages
	Browser	Zoom at 110-125% for readability; close all other tabs and notifications
	Audio	Microphone permissions granted; TTS output volume tested
	Recorder	Screen recorder ready (OBS / QuickTime); set to 1920x1080
	Sidebar	Expanded (not collapsed)
	Pre-test	Run one text query AND one mic query to confirm backend responds

RECORDING TIPS

- **Record screen first, add voiceover separately.** This drastically reduces retakes.
- **Mouse movements:** Smooth and deliberate. No jittery clicking.
- **Pacing:** Pause 2-3 seconds on each important visual before moving on.
- **Typing:** Type queries at a natural pace - not too fast, not hunting for keys.
- **Voice demo:** Speak clearly at a normal pace. Rehearse 2-3 times.
- **Captions:** Add subtitles in post-production. 80% of social video is watched on mute.
- **Export formats:** 1920x1080 (16:9) for Twitter. 1080x1080 (square) for LinkedIn.
- **Thumbnail:** Use a frame from the Command Center map view with facility markers.

POST-RECORDING VERIFICATION

- Watch the recording at 1x speed - should be between 5:00 and 6:00
- Ensure every screen is visible for at least 10 seconds
- Confirm the voice demo (mic input + TTS playback) is audible and clear
- Check that no sensitive data (API keys, .env files) is visible on screen
- Verify captions/subtitles are synced if added
- Test playback on mobile (small screen readability check)